

Introduction & an Overview of Relational Database Systems

CSCI 550: Advanced Data Stores
Spring 2024

Ibrahim Sabek
sabek@usc.edu

Course Staff



Ibrahim Sabek (Instructor)
sabek@usc.edu



Meet Raval (TA)
mkraval@usc.edu



Vash Singh (CP)
singhvas@usc.edu

Administrivia

Course materials

Instructor topics, CMU (Advanced DB Systems), DB Redbook (www.redbook.io/)

No textbook! Mostly research papers

Lectures will be posted on Blackboard

Ask questions on Piazza! Use signup link

Gradescope for uploading reports and grading! Use entry code

Instructor office hours

TuTh 4:00–5:00 pm

Office: PHE 420

TA office hours: TBD

Administrivia

No lecture recording, unless stated otherwise!

In-person attendance, unless permission is granted

Interactive class! Raise your hand if you have questions

2-hours class

1 hour for the instructor's materials

1 hour for the students' presentations (3 20-min presentations, individual)

Starting on Jan 16 (Tuesday)

15% of total grade

Guest lectures (maybe!)

Administrivia

Review reports (25% of total grade)

Conference review style!

8 reports (almost weekly, topic-based, individual)

1st report due is Jan 23 (Tuesday)

Late submission policy (max of 3 days late per report, reduce max. credit by 20% each day)

Checkpoint exams (30% of total grade)

2 exams (middle of the semester, end of the semester)

Course project (30% of total grade)

Groups of 2-3

Extend an open-source DBMS with a feature!

Project proposal, implementation, presentation, and conference paper-style report

1 status update meeting

Academic honesty

Refer to the [USC Academic Integrity Policy](#)

Plagiarism will **not** be tolerated (writing and implementation must be your own)

Why Should I Take CSCI 550?

Designing and building database management systems (DBMSs), not how to use them!

We will cover state-of-the-art and modern DBMS techniques (not classical DBMSs)

Mostly in-memory DBMSs (components + overall system design)

DBMS developers are needed everywhere!

What Is A Database?

Structured Data Collection (Tuples + Relationships)

Database Management Systems (DBMSs)

Software systems for storing, querying, and processing data in databases



Source: mattturck.com

Main Modern (In-memory) DBMS Concepts

Indexing (OLAP and OLTP)

Query Pipelining and Materialization

Query Scheduling

Vectorized Execution

Query Compilation and Optimization

Consistency and Concurrency Control

Scalable System Architectures

A Database from A User Perspective

How to model and query data as a user

Zoo website features

Admin interface

- Add and edit an animal

Public interface

- Pictures and maps

Zookeeper interface

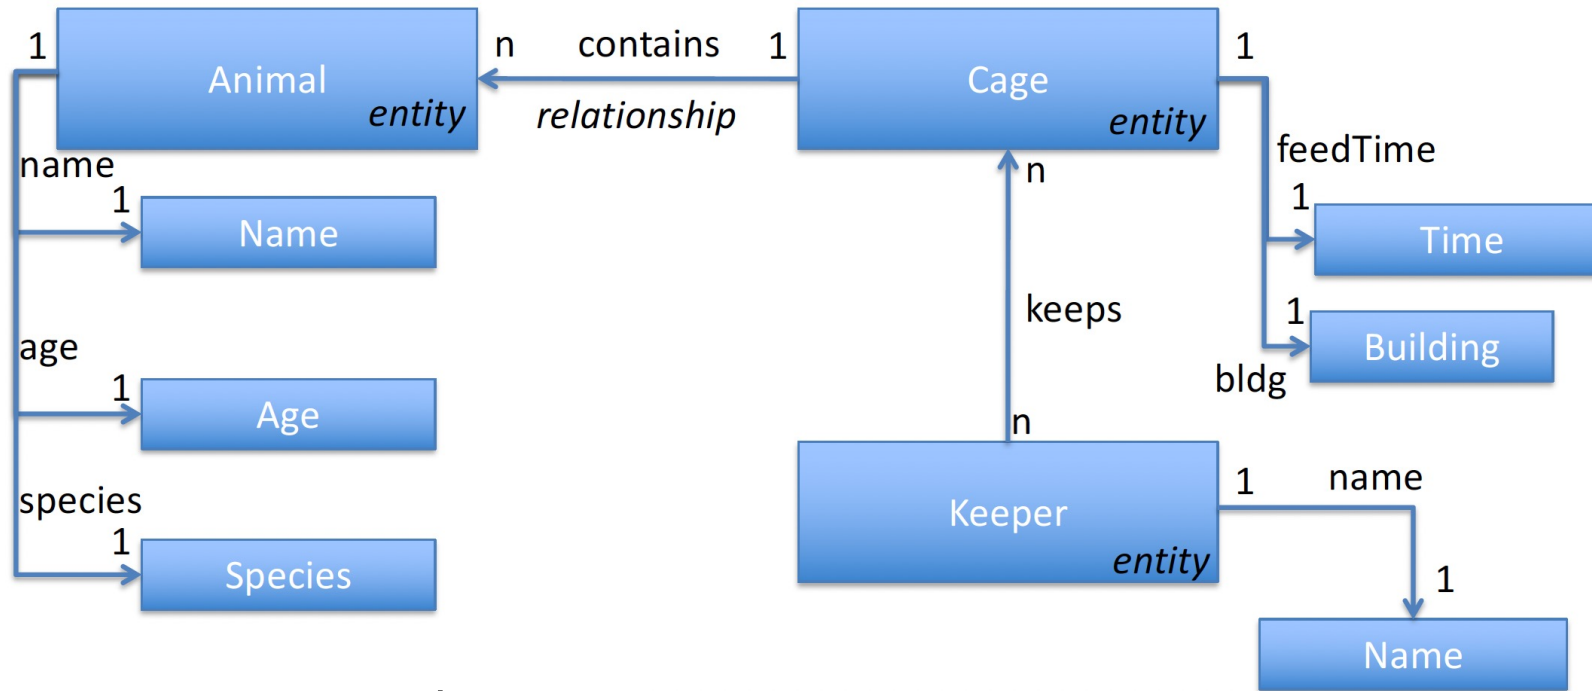
- Feed times

1K animals, 5K URLs, 10 admins, 200 keepers



Zoo Data Model

Entity Relationship Diagram (ERD)



Animals have names, ages, species

Keepers have names

Cages have cleaning times, and buildings

Animals are in 1 cage; cages have multiple animals

Keepers keep multiple cages, cages kept by multiple keepers

Structured Query Language (SQL)

```
SELECT  field1, ..., fieldM  
FROM    table1, ...  
WHERE   condition1, ...
```

```
INSERT INTO table VALUES (field1, ...)
```

```
UPDATE table SET field1 = X, ...  
WHERE condition1,...
```

Structured Query Language (SQL)

Names of giraffes

Imperative

```
for each row r in animals
    if r.species = 'giraffe'
        output r.name
```

Declarative

```
SELECT r.name FROM animals
WHERE r.species = 'giraffe'
```

Structured Query Language (SQL)

Cages in Building 32

Imperative

```
for each row a in animals
  for each row c in cages
    if a.cageno = c.no and c.bldg = 32
      output a
```

Nested Loop

Declarative

```
SELECT a.name FROM animals AS a, cages AS c
WHERE a.cageno = c.no AND c.bldg = 32
```

Join

Structured Query Language (SQL)

Average age of bears

Declarative

```
SELECT AVG(age) FROM animals  
WHERE species = 'bear'
```

Declarative Queries: What Not How!

Many possible procedural plans for a given SQL query

What else could we do for the "bear" query?

- Sort animals according to type

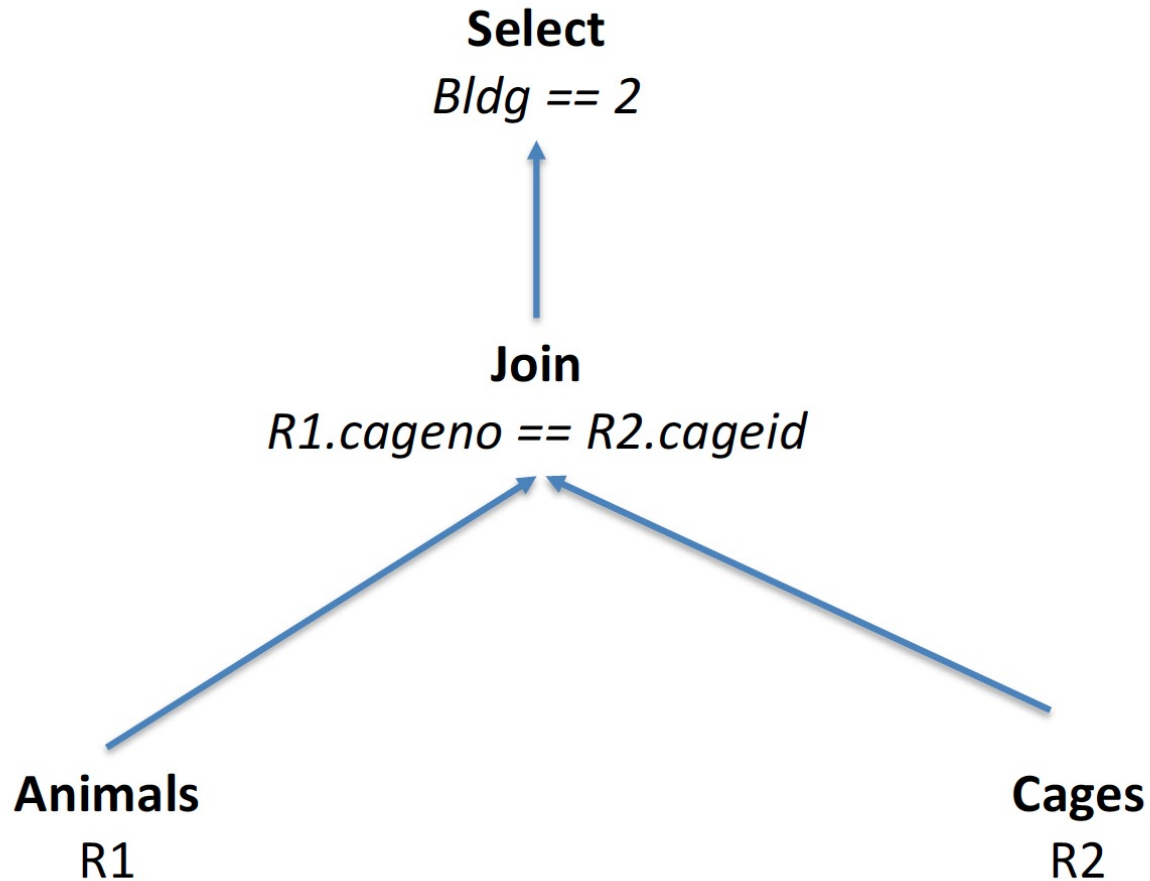
 - good for the "bear" query

 - inserts are slower

- Store animals table in a hash table or tree ("index")

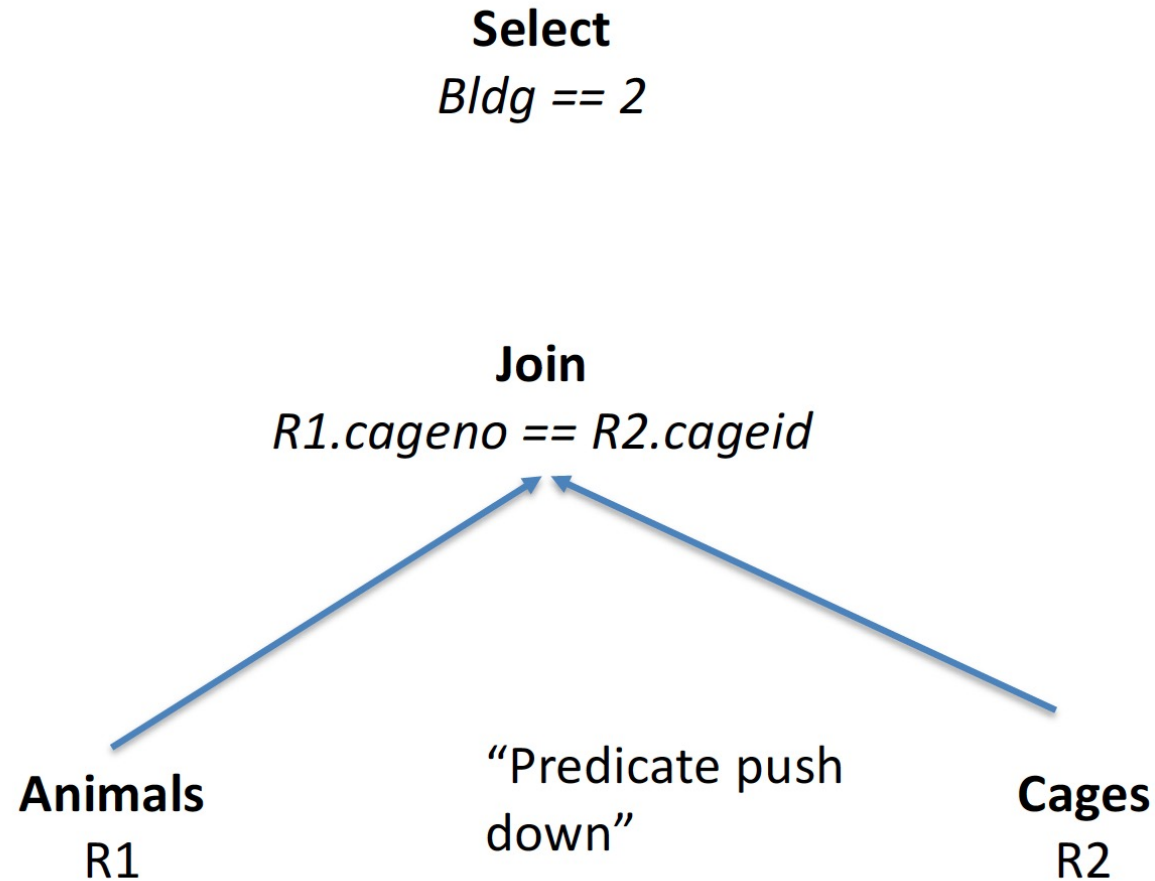
From Logical Plan to An Optimized Plan

Cages in Building 2



From Logical Plan to An Optimized Plan

Cages in Building 2



Summary

DBMSs are the key success to many revolutionary applications we witness

Different components interact together

- Data Modeling

- Declarative Querying

- Query Optimization

- Efficient access and updates to data

 - Recoverability

 - Consistency

Administrivia: be active and ask if you have questions